

Overview

Role
ML Engineer

Session time
15m 44s

Questions answered
8

Completion
100%

Metrics

Overall score

51.88

Q1

77/100

Q2

88/100

Q3

80/100

Q4

30/100

Q5

30/100

Q6

15/100

Q7

35/100

Q8

60/100


MIXED

HARD

COMPLETED

Interview Report

• Sep 14, 2026 • 40 min • 1 rounds



- Your design demonstrates a solid grasp of modern ANN techniques and distributed system considerations, presenting a clear end-to-end pipeline.
- Including more detail on index updates, fault tolerance, and latency budgeting would make the solution even stronger.
- Your answer presents several practical techniques for streaming and memory-efficient ANN search, such as quantization and a dual-buffer architecture.
- It would be stronger if you linked these changes more explicitly to the original algorithm's components and explained how incremental updates are handled.

Category performance

Technical	52.5%
Communication	66.25%
Problem solving	61.25%
Confidence	45.63%

Adaptive journey

INTERVIEW COMPLETED

THE ROUTER ADAPTED THE INTERVIEW 0 TIMES BASED ON YOUR PERFORMANCE

Strengths

- ✓ Effective combination of HNSW and Product Quantization
- ✓ Clear description of data ingestion and sharding
- ✓ Scalable architecture with horizontal scaling
- ✓ Clear, structured presentation of ideas
- ✓ Good awareness of memory constraints and latency
- ✓ Concrete engineering solutions (quantization, ring buffer)
- ✓ Mentions caching (Varnish)
- ✓ Suggests indexing on user_id
- ✓ Identifies need for score updates
- ✓ Clear low-latency goal
- ✓ Identified candidate generation and ranking stages
- ✓ Mentioned caching and lightweight model inference

Areas to improve

- Limited discussion of index maintenance and consistency
- Missing quantitative latency/throughput estimates
- No mention of monitoring or fault-tolerance mechanisms
- Limited connection to the specific algorithm discussed
- Some concepts introduced without detailed justification
- No explicit discussion of incremental index maintenance
- Oversimplified storage choice

Recommended resources TIED TO OBSERVED GAPS

LIMITED DISCUSSION OF INDEX MAINTENANCE AND CONSISTENCY

Review this area and practise a focused follow-up response.

MISSING QUANTITATIVE LATENCY/THROUGHPUT ESTIMATES

Review this area and practise a focused follow-up response.

NO MENTION OF MONITORING OR FAULT-TOLERANCE MECHANISMS

Review this area and practise a focused follow-up response.

LIMITED CONNECTION TO THE SPECIFIC ALGORITHM DISCUSSED

Review this area and practise a focused follow-up response.

SOME CONCEPTS INTRODUCED WITHOUT DETAILED JUSTIFICATION

Review this area and practise a focused follow-up response.

NO EXPLICIT DISCUSSION OF INCREMENTAL INDEX MAINTENANCE

Review this area and practise a focused follow-up response.

OVERSIMPLIFIED STORAGE CHOICE

Review this area and practise a focused follow-up response.

Question-level analysis

8 questions

Q1

MIXED

GOOD

Design a real-time recommendation system with low-latency inference

77/100

Q2

MIXED

STRONG

Design and implement a scalable nearest-neighbor retrieval system for high-dimensional vectors

88/100

Q3

MIXED

STRONG

Considering the algorithm you just described, how would you modify it to handle streaming data with limited memory?

80/100

Q4

MIXED

WEAK

Design a Real-Time Recommendation Serving System with Top-K Retrieval

30/100

Q5

MIXED

WEAK

Design a Scalable Real-Time Feature Store for Machine Learning

30/100

Q6

MIXED

WEAK

Design a Scalable Feature Store for Real-time ML Inference

15/100

Q7

MIXED

WEAK

Design a real-time feature retrieval system for online machine learning predictions

35/100

Q8

MIXED

GOOD

Design a Scalable Feature Store and Implement Approximate Nearest Neighbor Retrieval

60/100

History

Question 1

Design a real-time recommendation system with low-latency inference

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 2

Design and implement a scalable nearest-neighbor retrieval system for high-dimensional vectors

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 3

Considering the algorithm you just described, how would you modify it to handle streaming data with limited memory?

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 4

Design a Real-Time Recommendation Serving System with Top-K Retrieval

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 5

Design a Scalable Real-Time Feature Store for Machine Learning

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 6

Design a Scalable Feature Store for Real-time ML Inference

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 7

Design a real-time feature retrieval system for online machine learning predictions

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 8

Design a Scalable Feature Store and Implement Approximate Nearest Neighbor Retrieval

Answer submitted

MIXED response recorded.

Answer evaluated

AI evaluation completed.

Question 9

Design a scalable real-time recommendation service